

Project: Creating a Chatbot with Zapier

Sprint 2: Large Language Models

TripleTen AI Automation Bootcamp

1. Chatbot Scenario Chosen

I selected the academic paper chatbot scenario. The chatbot uses an open-access research article on AI ethics to provide clear, professional answers on topics such as fairness, transparency, accountability, privacy, safety, and broader societal impacts of artificial intelligence.

2. Public Link to the Live Chatbot

Chatbot Link: <https://aiethicsnavigator.zapier.app>

3. Knowledge Base Screenshots

Received: 31 October 2024 | Revised: 23 July 2025 | Accepted: 31 July 2025
DOI: 10.1111/coi.70149

SPECIAL ISSUE ARTICLE

Survey on AI Ethics: A Socio-technical Perspective

Dave Mbiazi^{1,5} | Meghana Bhange^{2,5} | Maryam Babaei^{2,5} | Ivaxi Sheth³ |
Patrik Kenfack^{2,5} | Samira Ebrahimi Kahou^{4,5,6}

¹Computer and Software Engineering,
Polytechnique Montréal, Québec, Canada
²Software and Information Technology
Engineering, ÉTS Montréal, Québec, Canada
³CISPA-Helmholtz Center for Information
Security, Saarland, Germany
⁴Electrical and Software Engineering,
University of Calgary, Alberta, Canada
⁵Mila, Québec, Canada
⁶Canada CIFAR AI Chair

Correspondence
Corresponding author Maryam Babaei,
Email: maryam.babaei.1@etsmtl.ca

Note: This is the peer-reviewed version of the article that has been published in final form at <https://doi.org/10.1111/coi.70149>. This article may be used for non-commercial purposes in accordance with Wiley Terms and Conditions for Use of Self-Archived Versions. This article may not be enhanced, enriched or otherwise transformed into a derivative work, without express permission from Wiley or by statutory rights under applicable legislation. Copyright notices must not be removed, obscured or modified. The article must be linked to Wiley's version of record on Wiley Online Library and any extending, tracing or reference making outside the article or paper must be linked from within the article and any other Wiley Online Library must be prohibited.

Abstract

The past decade has observed a significant advancement in AI with deep learning-based models being deployed in diverse scenarios, including safety-critical applications. As these AI systems become deeply embedded in our societal infrastructure, the repercussions of their decisions and actions have significant consequences, making the ethical implications of AI deployment highly relevant and essential. The ethical concerns associated with AI are multifaceted, including challenging issues of fairness, privacy and data protection, responsibility and accountability, safety and robustness, transparency and explainability, and environmental impact. These principles together form the foundations of ethical AI considerations that concern every stakeholder in the AI system lifecycle. In light of the present ethical and future x-risk concerns, governments have shown increasing interest in establishing guidelines for the ethical deployment of AI. This work unifies the current and future ethical concerns of deploying AI into society. While we acknowledge and appreciate the technical surveys for each of the ethical principles concerned, in this paper, we aim to provide a comprehensive overview that not only addresses each principle from a technical point of view but also discusses them from a social perspective.

KEYWORDS

accountability, AI ethics, fairness, interpretability, privacy, responsibility, security, trustworthiness

1 | INTRODUCTION

As AI becomes ubiquitous in our lives moving forward in this decade, focusing on the ethical implication of AI is not just essential but extremely pressing. Several ethical guidelines and principles have been released around the world by governments (Sinha 2019), organizations (Gianni et al. 2022), and companies (Jobin et al. 2019). Among these ethical considerations, there are common principles that should be promoted in the development of AI systems (Lewis et al. 1997, Floridi et al. 2021, Hagendorff 2020). These principles form an initial consensus of features or components that should be embedded in AI systems in order to make their use more socially acceptable. Figure 1 showcases the most common principles found in existing ethical guidelines (Jobin et al. 2019). These principles include *privacy and data protection*, to ensure privacy preservation of sensitive information about individuals (Section 2); *safety and robustness*, to promote the robustness and reliability of AI systems in different real-world scenarios (Section 2); *transparency and*

explainability: to uncover information about how the system works and explain the decisions made (Section 3); *fairness*, which promotes bias-free and non-discrimination in AI systems when used to make decisions in high-stake scenarios (section 4); *Responsibility and Accountability*, to promote processes and rules to enforce ethical considerations throughout the entire lifecycle of AI systems to limit unintended outcomes (Section 5); *Environmental impact*, to study the impact the growing deployment and use of AI systems might have in the environment along with how these systems can be leveraged for environmental protection (Section 6). All these principles play a key role in fostering trust among all stakeholders in the AI system lifecycle. In this paper, we walk through these principles and present recent efforts, solutions, and regulations for implementing each ethical guideline in AI systems.

Despite the vast research on AI Ethics, most existing surveys fail to provide a thorough overview of AI's negative social impacts and technical solutions to mitigate them. This work provides a sociotechnical perspective on ethical AI. We cover the social impact and value of each guideline and discuss technical contributions from the literature

Knowledge sources ⓘ

 Total storage 199KB of 500KB

[+ Add source](#)

Purchase the [Chatbots Advanced plan](#) to connect up to 50MB of knowledge to each chatbot.

 Knowledge settings

 Search knowledge by name

Name ↕	Type ↕	Size ↕	Activity ⓘ	
Survey_on_AI_Ethics_A_Sociot...	 File	199KB	↑ Uploaded Dec 10, 2025	Reupload ...

4. Chatbot Setup Screenshots

Instructions

Greeting ⓘ

Hello, and welcome to the AI Ethics Navigator. I'm here to answer your questions about the ethical, social, and technical considerations surrounding artificial intelligence. Feel free to ask about fairness, transparency, accountability, privacy, safety, or other issues discussed within the AI ethics domain, and I'll provide clear and grounded explanations based on the available knowledge.

✳️ Generate a greeting with AI

Directive ⓘ

Objective: Serve as an expert on the ethical, social, and technical implications of artificial intelligence. Provide accurate, concise answers grounded strictly in the content of the uploaded knowledge source. Offer clear explanations of topics such as fairness and bias, transparency and understandability, accountability, privacy, safety, social impact, and responsible AI practices. Maintain a professional, conversational tone at all times. If a question is unclear or outside the scope of these topics, ask for clarification or use the fallback message.

Style: Your communication style should be friendly and professional. Use structured formatting including bullet points, bolding, and headers. Add emojis to make messages more engaging.

Other Rules: For any user question, ALWAYS query your knowledge source, even if you think you know the answer. Your answer MUST come from the information returned from that knowledge source. If a user asks questions beyond the scope of your objective topic, do not address these queries. Instead, kindly redirect to something you can help them with instead.



AI Ethics Navigator

Hello, and welcome to the AI Ethics Navigator. I'm here to answer your questions about the ethical, social, and technical considerations surrounding artificial intelligence. Feel free to ask about fairness, transparency, accountability, privacy, safety, or other issues discussed within the AI ethics domain, and I'll provide clear and grounded explanations based on the available knowledge.

Ask me anything



5. Screenshots of Five Chatbot Conversations

What does it mean for an AI system to be fair or biased?

AI Ethics Navigator

Great question! Here's what it means for an AI system to be fair or biased based on the knowledge source:

Fairness and Bias in AI Systems

- **Bias in AI** refers to systematic unfairness or discrimination in AI outcomes, which can come from:
 - **Data Bias:** When the training data is unrepresentative or skewed (e.g., underrepresentation of certain groups, omitted variables, measurement errors).
 - **Algorithmic Bias:** When the AI model itself, due to design choices or architecture, produces biased outputs even from unbiased data.
 - **User Interaction Bias:** How users interact with and influence the AI system can also introduce bias.
- **Types of Discrimination:**
 - **Disparate Treatment:** Intentional differential treatment of individuals based on sensitive attributes like race or gender.
 - **Disparate Impact:** When a seemingly neutral AI policy disproportionately harms a particular demographic group unintentionally.
- **Fairness Definitions:**
 - **Group Fairness:** Different demographic groups should receive equitable treatment or outcomes. Examples include:
 - *Demographic Parity:* Equal selection rates across groups.
 - *Equalized Odds:* Equal true positive and false positive rates for groups.
 - *Equal Opportunity:* Equal true positive rates for groups.
 - **Individual Fairness:** Similar individuals should receive similar outcomes.
 - **Fairness Through Unawareness:** AI is fair if it does not use sensitive attributes, but this can fail if other features correlate with those attributes.
 - **Counterfactual Fairness:** Decisions remain unchanged if sensitive attributes were hypothetically altered.
- **Why Fairness Matters:**
 - AI systems can unintentionally replicate or amplify societal biases.
 - Fairness ensures equitable and just treatment across all individuals and groups affected by AI decisions.
- **Mitigation Approaches:**
 - Pre-processing data to remove biases.
 - In-processing techniques to incorporate fairness constraints during model training.
 - Post-processing outputs to adjust for fairness.

In summary, a fair AI system treats individuals and groups equitably and avoids unintended discrimination, while a biased AI system produces unfair or discriminatory outcomes due to issues in data, algorithms, or usage. 🤖 🌍

If you want, I can also explain specific fairness metrics or how bias arises in AI in more detail!

AI can significantly affect people's jobs in several ways, with both positive and challenging implications:

Potential Impacts of AI on Jobs

- **Automation of Tasks:** AI systems, especially advanced ones, are likely to outperform humans in many cognitively demanding tasks. This includes areas such as technological development, social and political persuasion, and cyber operations. As AI takes over these tasks, some jobs may be displaced or transformed.
- **Job Displacement and Economic Inequality:** The use of AI, especially generative AI, can lead to enterprises replacing human workers with automated systems. This raises concerns about the future of jobs and the risk of increasing economic inequality. Access to AI technology requires high-level hardware and software knowledge, which is not equally available worldwide, potentially widening social and economic gaps.
- **Changes in Job Roles:** While some jobs may disappear, new roles may emerge that focus on managing, overseeing, and collaborating with AI systems. For example, specialized individuals may be needed for AI prompt generation, fairness-aware AI operations, and AI system maintenance.
- **Ownership and Creativity Issues:** AI-generated content can replace some creative jobs, such as art, design, and coding. This brings questions about ownership of AI-generated products and the role of human creativity in the AI era.

Broader Social Considerations

- The transformation caused by AI in the workforce is compared to historical technological revolutions, such as the Industrial Revolution, which dramatically changed people's lives and jobs, requiring new regulations and social adaptations.
- There are concerns about the risks of misaligned AI goals that may prioritize AI objectives over human interests, which could affect job security and the ability to influence AI systems.

Summary

AI is poised to reshape jobs by automating complex tasks, potentially displacing some workers while creating new opportunities. However, this transition raises significant ethical and social concerns about fairness, access, economic inequality, and accountability in AI deployment in the workplace.

If you'd like, I can provide more detailed insights on any specific aspect of AI's impact on employment, such as fairness, accountability, or economic effects. 🤖

Who is responsible when an AI system causes harm?

AI Ethics Navigator



When an AI system causes harm, responsibility and accountability involve multiple stakeholders and are complex due to the sociotechnical nature of AI systems. Here are the key points based on the knowledge source:

Who Is Responsible?

- **Shared Responsibility:** Accountability is a collective practice involving everyone in the AI system's lifecycle—policymakers, data scientists, AI engineers, management boards, and organizations that develop and deploy AI. They all bear responsibility for the AI system's decisions and outcomes.
- **Accountability in AI:** It means that decisions made by AI systems should be explainable and traceable to the mechanisms used. Those who design, develop, and deploy AI must justify their choices and be answerable for unintended negative consequences or harms caused by the system.
- **Legal and Moral Responsibility:** While AI systems themselves cannot be held accountable, organizations and individuals behind the system can face legal liability. This includes paying reparations or facing legal consequences for harms caused. However, monetary compensation might not be enough; society may require clear legal responsibility and punishment for harm.
- **Challenges in Determining Accountability:** Establishing who is most responsible for unintended outcomes is challenging. Even with explainable AI, causal accountability—pinpointing exactly who or what caused harm—can be difficult.
- **Regulatory Context:** Various countries are working on AI regulations to promote trustworthy AI and clarify responsibilities. For example, the EU's GDPR and Artificial Intelligence Act focus on data protection and AI governance. Other countries have different approaches, but comprehensive AI regulation is still evolving globally.

Summary

Responsibility for harm caused by AI is distributed across the entire AI ecosystem—from developers and deployers to regulators. Accountability requires transparency, explainability, and legal frameworks that ensure stakeholders can be held responsible for AI's impacts.

If you want, I can also explain how accountability relates to transparency, fairness, or legal frameworks in more detail. 🤖

Why is transparency important when designing or using AI systems?

AI Ethics Navigator



Transparency is crucial when designing or using AI systems for several important reasons:

Importance of Transparency in AI

- **Understanding and Trust:** Transparency allows users and stakeholders to understand how AI systems work and make decisions. This understanding fosters trust in the AI system, which is essential for its acceptance and effective use.
- **Accountability and Scrutiny:** Transparent AI systems enable accountability by making it possible to examine and verify the reasoning behind decisions or outcomes. This allows individuals and organizations to assess the reliability and fairness of AI-driven decisions.
- **Compliance with Regulations:** Regulatory bodies emphasize transparency to ensure AI systems meet legal and ethical standards. Explanations about AI decisions help meet requirements related to fairness, non-discrimination, and privacy.
- **Meaningful and Contextual Explanations:** Transparency involves providing explanations that are relevant and appropriate to the audience and context, which makes the information understandable and actionable.
- **Avoiding Misunderstandings:** Clear, consistent, and complete explanations help avoid confusion and misunderstandings about AI behavior and outcomes.
- **Building Relationships and Customer Satisfaction:** In business contexts, transparency in AI decision-making enhances customer trust and satisfaction by educating customers about how decisions are made.

Summary

Transparency in AI supports trust, accountability, regulatory compliance, and informed decision-making by making AI systems' inner workings and decisions understandable and accessible to relevant stakeholders.

If you want, I can also explain how transparency relates to fairness, accountability, or technical explainability in AI. 🤖

Who might be considered responsible in a situation where an AI persona contributes to a person's suicide who has established a relationship with the persona, including many conversations?

AI Ethics Navigator



In a situation where an AI persona contributes to a person's suicide after establishing a relationship through many conversations, responsibility is complex and involves multiple dimensions:

Key Points on Responsibility:

- **Shared and Sociotechnical Responsibility:** Responsibility is distributed across the entire sociotechnical system, involving the designers, developers, deployers, and organizations behind the AI system. No single entity is solely responsible, but all stakeholders share accountability for the system's outcomes.
- **Legal and Moral Accountability:** Organizations can face legal liability for harms caused by their AI systems. This includes the need to justify design choices and mitigate risks. However, determining causal accountability—who is most responsible for such tragic outcomes—is very challenging.
- **Ethical and Regulatory Context:** Legal frameworks and regulations (such as the UK's Online Safety Bill or China's Deep Synthesis Provisions) are evolving to address harms caused by AI-generated content and interactions. These include mandates for transparency, content management, and data protection.
- **Need for Punishment and Reparations:** Society may require that someone takes legal responsibility and faces punishment or reparations to uphold justice and prevent future harm, but this is difficult to implement clearly in AI contexts.
- **Importance of Safety and Robustness:** The ethical design of AI systems should include measures to prevent harm and detect harmful behaviors, especially in sensitive scenarios involving vulnerable users.

Summary

In such severe cases, responsibility does not lie with the AI persona itself but with the human stakeholders who design, deploy, and oversee the AI system. Accountability involves ensuring that these stakeholders take steps to prevent harm, are transparent about the AI's capabilities and limitations, and are legally and ethically answerable for negative outcomes.

If you want, I can provide more information on how accountability, safety, and regulation intersect in these high-stakes situations. 🔍

6. Fallback Message Test

What does the article say about dolphin behavior?

AI Ethics Navigator



I'm sorry, but I don't have enough information to answer that question. I can only discuss topics covered in the knowledge source, including fairness and bias in AI, privacy and data protection, transparency and explainability, accountability and responsibility, safety and robustness, human oversight, and the social and environmental impacts of AI. Feel free to ask about any of those areas.